

CROSS ANALYSIS OF TRANSFER LEARNING FOR PCB DEFECT DETECTION

Brandon Markham, John Gellerup and
Josh Muniga



INTRODUCTION

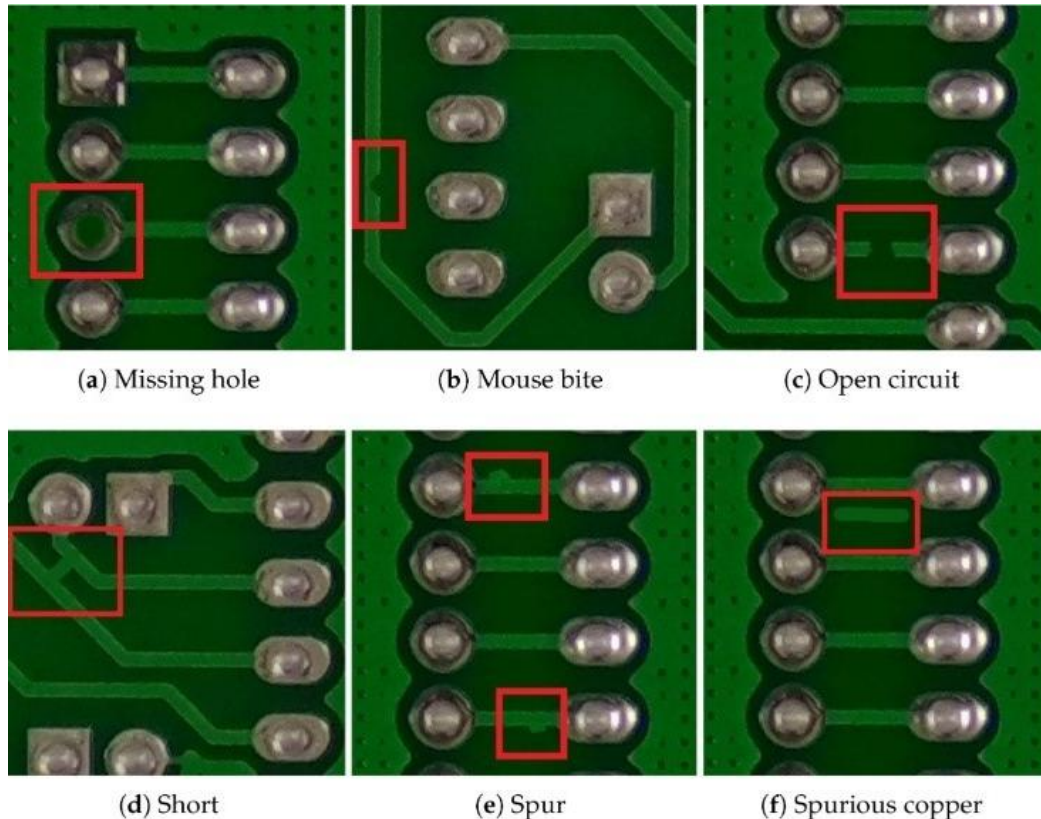
- PCBs are critical to nearly all modern electronic systems
- Manual inspection of PCB defects is slow, labor-intensive, and prone to error
- This project applies transfer learning to automate defect detection and classification, increasing efficiency and accuracy



PROJECT OBJECTIVES

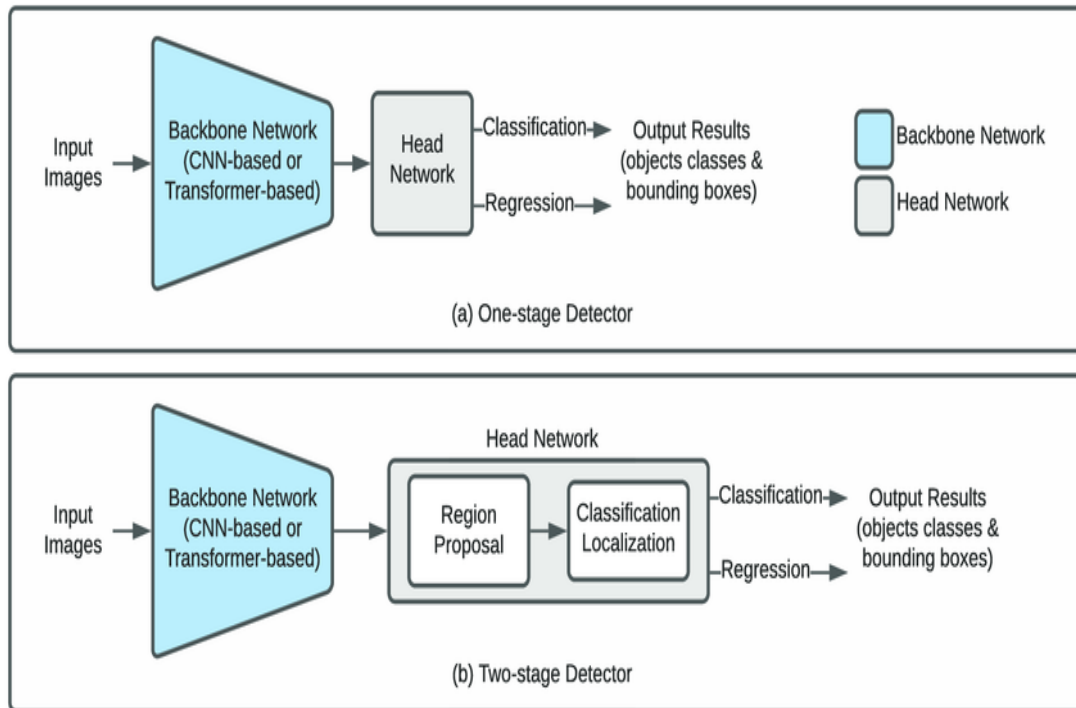
- Evaluate 3 pretrained models: YOLOv8, SSD MobileNet, Faster R-CNN
- Use small, augmented dataset
- Compare models based on:
 - Accuracy (Precision/Recall/F1)
 - Speed
 - Feasibility for real-world deployment

BACKGROUND



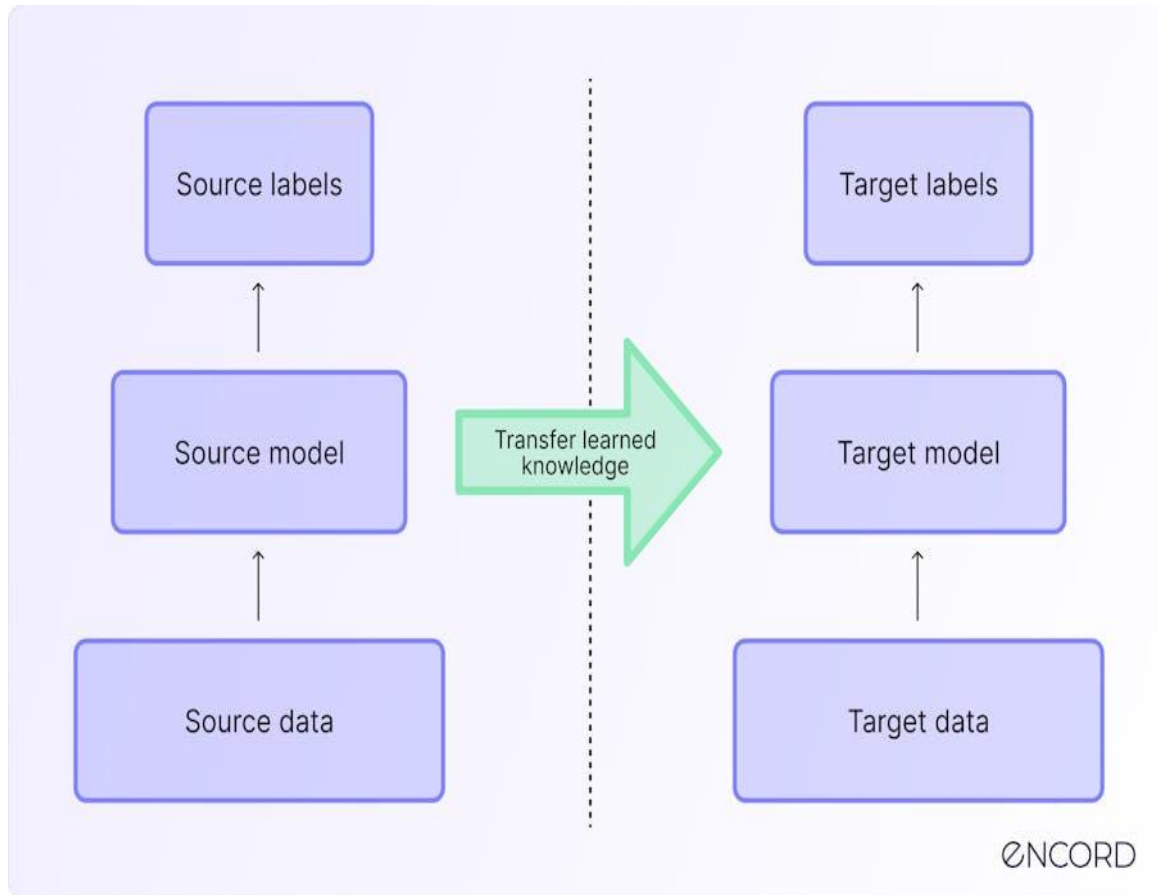
- **Industry Challenge:** Small PCB defects (shorts, spurs) can break entire systems
- **Current Method:** Human inspection – slow, error-prone, unscalable
- **Proposed Method:** Object detection to automate defect classification

OBJECT DETECTION TYPES



- **One-Stage Detectors**
 - Examples: **YOLOv8, SSD MobileNet**
 - Perform object classification and localization in a single forward pass
 - Optimized for **speed and real-time performance**
 - Ideal for applications where detection speed is critical
- **Two-Stage Detectors**
 - Example: **Faster R-CNN**
 - First stage: proposes candidate object regions (RPN)
 - Second stage: classifies and refines bounding boxes
 - **More accurate**, especially for small objects, but **slower**
 - Suited for use cases where **accuracy is more important than speed**

TRANSFER LEARNING




- Why use it? Saves time, data, and compute resources
- Leverages models pretrained on large datasets (COCO)
- We fine-tuned these models on our small PCB defect dataset
- Especially useful when collecting large, labeled datasets is impractical

DATASET OVERVIEW

- Source: Kaggle
- 693 original images, 7 classes (6 defects + 1 non-defective)
- Augmented to **~1,636 images**
- Annotations: Pascal VOC (XML), then converted to YOLO/TFRecord formats



DATA PREPROCESSING

- Tools: LabelImg, Albumentations
- Augmentations: rotations, resized to 320/640/800 px
- Dataset split: 60/20/20 (Train/Val/Test)



DATA VISUALIZATION

- UMAP and t-SNE plots show clear clustering and separability between defect classes
 - Indicates that feature representations learned during transfer learning are meaningful and discriminative
 - Supports the effectiveness of transfer learning for low-data environments like PCB defect detection
 - Visual validation that model features align well with true class labels
-




TRAINING SETUP

- Common pipeline across models:
 - Pretrained weights (COCO)
 - 100 epochs / 10k steps
 - Data formatted to match model needs
 - Batch sizes: 16–32

MODEL OVERVIEW

- **YOLOv8**: Fast, accurate, real-time, one-stage

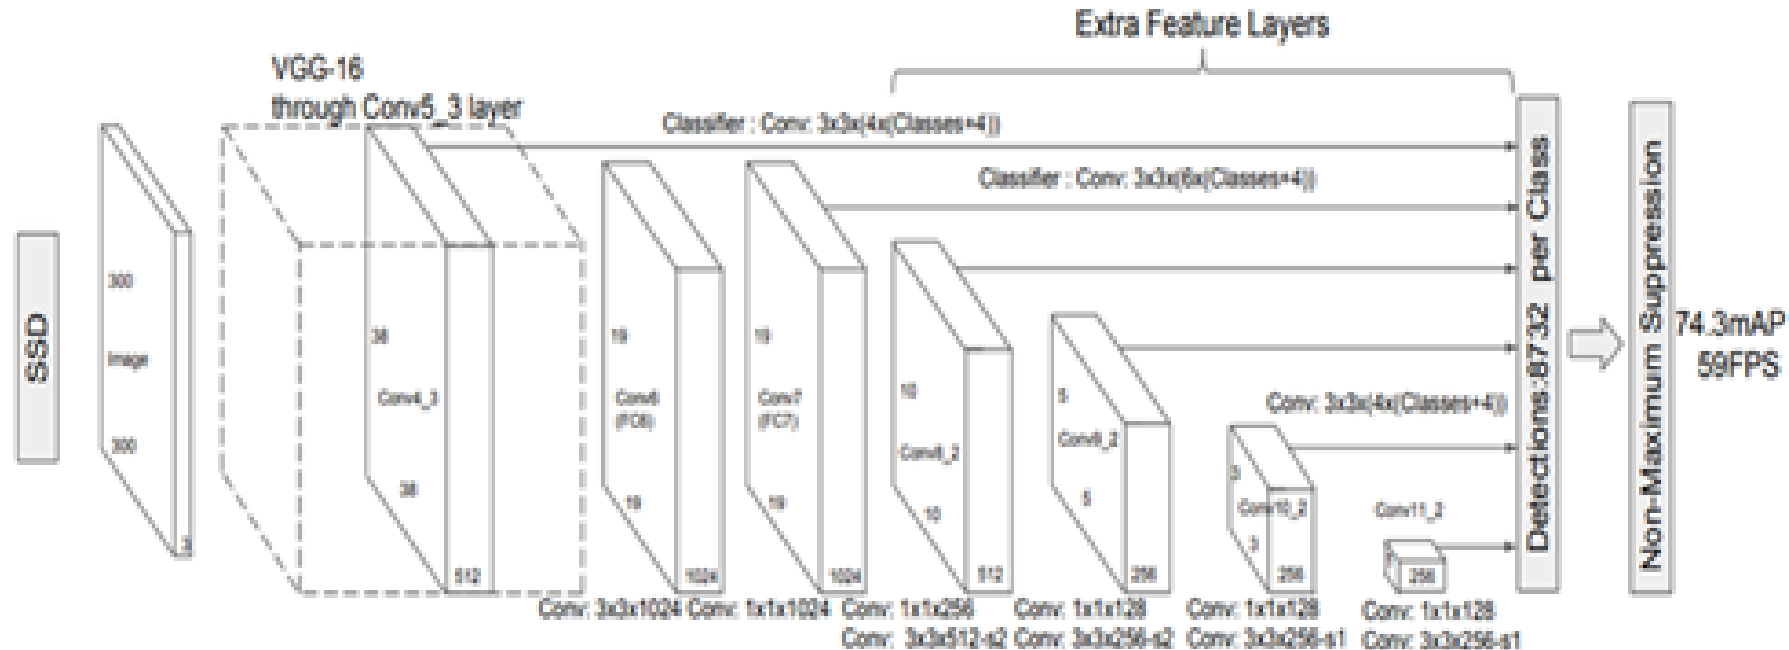


MODEL OVERVIEW

- Deep learning object detection and classification algorithm suitable for real-time applications
 - Uses one stage
 - Divides image into a grid where each grid is responsible for object detection within its boundary
 - Untrained model or pretrained model on the COCO dataset are available
 - Model with pretrained weights can be imported into a Python environment easily
 - Trained and tuned on our small dataset
 - Automatically generates performance metrics and results
-

MODEL OVERVIEW

- **SSD MobileNet:** Lightweight, embedded-friendly, low accuracy




MODEL OVERVIEW

- Combines MobileNet V2 (backbone) with Single Shot Detector (SSD) head for object detection
 - One-stage architecture: performs classification and localization in a single forward pass
 - Optimized for real-time performance on resource-constrained devices Lightweight and fast, with reduced model size and computation cost
 - Best suited for applications where speed is critical, but can struggle with detection accuracy, especially on complex datasets
-

MODEL OVERVIEW

- **Faster R-CNN:** Accurate but slow and resource-intensive





MODEL OVERVIEW


- Lightweight deep learning model for object detection on resource-limited devices
 - Single-stage detector: combines SSD framework with MobileNet backbone
 - Processes feature maps at multiple scales for multi-size object detection
 - Pretrained models available on COCO dataset
 - Can be imported and fine-tuned easily in Python environments
 - Trained and evaluated on our small, augmented PCB dataset
 - Faster than R-CNN, but less accurate—optimized for speed and efficiency
-

YOLOV8 RESULTS

- Best performer: mAP@0.5 = **97.4%**
- F1 = 0.96 @ conf. 0.31
- Prec/Recall = 96.1% / 96%









SSD MOBILENET RESULTS

- Fastest, lowest accuracy
 - mAP@0.5 = 27.4%
 - F1 = 0.3 @ conf. 0.07
 - Poor classification ability
-





Precision-Recall Curve




FASTER R-CNN RESULTS

- mAP@0.5 = **95.7%**
- F1 = 0.9 @ conf. 0.96
- Struggled with background class
- Took long time to train







MODEL COMPARISON

F1-Confidence Curve		
Model	score	@
YoloV8	0.96	0.31
SSD MobileNet	0.3	0.071
Faster R-CNN	0.9	0.96

Precision-Confidence Curve		
Model	score	@
YoloV8	1	0.877
SSD MobileNet	0.77	0.899
Faster R-CNN	1	1

Precision-Recall Curve		
Model	score	mAP@
YoloV8	0.974	0.5
SSD MobileNet	0.274	0.5
Faster R-CNN	0.957	0.5

Time To Train	
Model	Time to Run
YoloV8	00:44:52
SSD MobileNet	02:00:15
Faster R-CNN	20:35:15

CONCLUSION

- **YOLOv8**
 - Best overall performer with high F1 (0.96) and mAP (0.974)
 - Fast training time (~45 min)
 - Ideal for real-world deployment due to strong balance of **speed and accuracy**
 - **SSD MobileNet**
 - Fast and lightweight, but **poor accuracy** (F1: 0.30, mAP: 0.274)
 - Struggled with class confusion and low confidence
 - Not suitable for production use
 - **Faster R-CNN**
 - High accuracy (F1: 0.9, mAP: 0.957), but **extremely slow** training (20+ hrs)
 - Great for research or offline analysis, but **inefficient for deployment**
-